

Fine-Tuning GPT-4.1-mini

Copyright 2025 Denis Rothman

This notebook upgrades fine-tuning to GPT 4.1.

[OpenAI fine-tuning documentation](#)


Check the cost of fine-tuning your dataset on OpenAI before running the notebook.

Run this notebook cell by cell to:

1. Download and prepare the SQuAD dataset Stanford Question Answering Dataset (SQuAD) is a reading comprehension dataset.
2. Fine-tune a model
3. Run a fine-tuned model

✓ Installing the environment


```
1 #You can retrieve your API key from a file(1)
2 # or enter it manually(2)
3 #Comment this cell if you want to enter your key manually.
4 #(1)Retrieve the API Key from a file
5 #Store you key in a file and read it(you can type it directly in the notebook but it will be visible for somebody next
6 from google.colab import drive
7 drive.mount('/content/drive')
8 f = open("drive/MyDrive/files/api_key.txt", "r")
9 API_KEY=f.readline()
10 f.close()
```

 Mounted at /content/drive

```
1 try:
2     import openai
3 except:
4     !pip install openai==1.42.0
5     import openai
```

```
1 #(2) Enter your manually by
2 # replacing API_KEY by your key.
3 #The OpenAI Key
4 import os
5 os.environ['OPENAI_API_KEY'] =API_KEY
6 openai.api_key = os.getenv("OPENAI_API_KEY")
```

```
1 !pip install jsonlines==4.0.0
```

 Collecting jsonlines==4.0.0
 Downloading jsonlines-4.0.0-py3-none-any.whl.metadata (1.6 kB)
Requirement already satisfied: attrs>=19.2.0 in /usr/local/lib/python3.11/dist-packages (from jsonlines==4.0.0) (25.3.0)
 Downloading jsonlines-4.0.0-py3-none-any.whl (8.7 kB)
Installing collected packages: jsonlines
Successfully installed jsonlines-4.0.0

```
1 !pip install datasets==2.20.0
```

 [Show hidden output](#)

Listing the installed packages

```
1 import subprocess
2
3 # Run pip list and capture the output
4 result = subprocess.run(['pip', 'list'], stdout=subprocess.PIPE, text=True)
5
6 # Split the output into lines and count them
```

```

7 package_list = result.stdout.split('\n')
8
9 # Adjust count for headers or empty lines
10 package_count = len([line for line in package_list if line.strip() != '']) - 2
11
12 print(f"Number of installed packages: {package_count}")

```

➞ Number of installed packages: 636

```

1 import subprocess
2
3 # Run pip list and capture the output
4 result = subprocess.run(['pip', 'list'], stdout=subprocess.PIPE, text=True)
5
6 # Print the output
7 print(result.stdout)

```

➞

google-pasta	0.2.0
google-resumable-media	2.7.2
googleapis-common-protos	1.70.0
googledrivedownloader	1.1.0
gradio	5.41.0
gradio_client	1.11.0
graphviz	0.21
greenlet	3.2.3
groovy	0.1.2
grpc-google-iam-v1	0.14.2
grpc-interceptor	0.15.4
grpcio	1.74.0
grpcio-status	1.71.2
grpccli	0.4.8
gsread	6.2.1
gsread-dataframe	4.0.0
gym	0.25.2
gym-notices	0.1.0
gymnasium	1.2.0
h11	0.16.0
h2	4.2.0
h5netcdf	1.6.4
h5py	3.14.0
hdbscan	0.8.40
hf_transfer	0.1.9
hf-xet	1.1.7
highspy	1.11.0
holidays	0.78
holoviews	1.21.0
hpack	4.1.0
html5lib	1.1
httpcore	1.0.9
httpimport	1.4.1
httplib2	0.22.0
httpx	0.28.1
huggingface-hub	0.34.3
humanize	4.12.3
hyperframe	6.1.0
hyperopt	0.2.7
ibis-framework	9.5.0
idna	3.10
imageio	2.37.0
imageio-ffmpeg	0.6.0
imagesize	1.4.1
imbalanced-learn	0.13.0
immutabledict	4.2.1
importlib_metadata	8.7.0
importlib_resources	6.5.2
imutils	0.5.4
inflect	7.5.0
iniconfig	2.1.0
intel-cmplr-lib-ur	2025.2.0
intel-openmp	2025.2.0
ipyevents	2.0.2
ipyfilechooser	0.6.0
ipykernel	6.17.1
ipyleaflet	0.20.0
ipyparallel	8.8.0
ipython	7.34.0

counting the number of packages

1.Preparing the dataset for fine-tuning

1.1.Downloading and displaying the dataset

```
1 from datasets import load_dataset
2 import pandas as pd
3
4 # Load the SQuAD dataset from HuggingFace
5 dataset = load_dataset("squad", split="train[:500]")
6
7 # Filter the dataset to ensure context and answer are present
8 filtered_dataset = dataset.filter(lambda x: x["context"] != "" and x["answers"]["text"] != [])
9
10 # Extract prompt (context + question) and response (answer)
11 def extract_prompt_response(example):
12     return {
13         "prompt": example["context"] + " " + example["question"],
14         "response": example["answers"]["text"][0] # Take the first answer
15     }
16
17 filtered_dataset = filtered_dataset.map(extract_prompt_response)
18
19 # Print the number of examples
20 print("Number of examples: ", len(filtered_dataset))
```

Show hidden output

```
1 # Convert the filtered dataset to a pandas DataFrame
2 df_view = pd.DataFrame(filtered_dataset)
3
4 # Display the DataFrame
5 df_view.head()
```

	id	title	context	question	answers	prompt	response
0	5733be284776f41900661182	University_of_Notre_Dame	Architecturally, the school has a Catholic cha...	To whom did the Virgin Mary allegedly appear i...	{'text': ['Saint Bernadette Soubirous'], 'answ...	Architecturally, the school has a Catholic cha...	Saint Bernadette Soubirous
1	5733be284776f4190066117f	University_of_Notre_Dame	Architecturally, the school has a Catholic cha...	What is in front of the Notre Dame Main Building?	{'text': ['a copper statue of Christ'], 'answe...	Architecturally, the school has a Catholic cha...	a copper statue of Christ
2	5733be284776f41900661180	University_of_Notre_Dame	Architecturally, the school has a Catholic cha...	The Basilica of the Sacred heart at Notre Dame...	{'text': ['the Main Building'], 'answer_start'...	Architecturally, the school has a Catholic cha...	the Main Building
3	5733be284776f41900661181	University_of_Notre_Dame	Architecturally, the school has a Catholic cha...	What is the Grotto at Notre Dame?	{'text': ['a Marian place of prayer and reflec...	Architecturally, the school has a Catholic cha...	a Marian place of prayer and reflection
			Architecturally,	What sits	{'text': ['a	Architecturally,	a golden

1.2A Streaming the output to JSON

```
1 import json
2 import pandas as pd
```

✓ 1.2. Preparing the dataset for fine-tuning

```
1 import jsonlines
2 import pandas as pd
3 from datasets import load_dataset
4
5 # Convert to DataFrame and clean
6 df = pd.DataFrame(filtered_dataset)
7 #columns_to_drop = ['title','question','answers']
8 #df = df.drop(columns=columns_to_drop)
9
10 # Prepare the data items for JSON lines file
11 items = []
12 for idx, row in df.iterrows():
13     detailed_answer = row['response'] + " Explanation: " + row['context']
14     items.append({
15         "messages": [
16             {"role": "system", "content": "Given a SQuAD question built from Wikipedia with crowdworkers, provide the"},
17             {"role": "user", "content": row['question']},
18             {"role": "assistant", "content": detailed_answer}
19         ]
20     })
21
22 # Write to JSON lines file
23 with jsonlines.open('/content/QA_prompts_and_completions.json', 'w') as writer:
24     writer.write_all(items)
```

✓ Visualizing the JSON file

```
1 dfile="/content/QA_prompts_and_completions.json"
```

```
1 import pandas as pd
2
3 # Load the data
4 df = pd.read_json(dfile, lines=True)
5 df
```



messages

0	[{"role": "system", "content": "Given a SQuAD ...
1	[{"role": "system", "content": "Given a SQuAD ...
2	[{"role": "system", "content": "Given a SQuAD ...
3	[{"role": "system", "content": "Given a SQuAD ...
4	[{"role": "system", "content": "Given a SQuAD ...
...	...
495	[{"role": "system", "content": "Given a SQuAD ...
496	[{"role": "system", "content": "Given a SQuAD ...
497	[{"role": "system", "content": "Given a SQuAD ...
498	[{"role": "system", "content": "Given a SQuAD ...
499	[{"role": "system", "content": "Given a SQuAD ...

500 rows × 1 columns

✓ 2. Fine-tuning the model

```
1 from openai import OpenAI
2 import jsonlines
3 client = OpenAI()
4 # Uploading the training file
5
6 result_file = client.files.create(
```

```

7 file=open("QA_prompts_and_completions.json", "rb"),
8 purpose="fine-tune"
9 )
10
11 print(result_file)
12 param_training_file_name = result_file.id
13 print(param_training_file_name)
14
15 # Creating the fine-tuning job
16 ft_job = client.fine_tuning.jobs.create(
17     training_file=param_training_file_name,
18     model="gpt-4.1-mini-2025-04-14"
19 )
20
21 # Printing the fine-tuning job
22 print(ft_job)

```

```

FileObject(id='file-6VgWqp3wmd6Tos6aUihLLZ', bytes=645098, created_at=1755077028, filename='QA_prompts_and_completions
file-6VgWqp3wmd6Tos6aUihLLZ
FineTuningJob(id='ftjob-f9JxcdrNVYiVn04M3EZkLaLT', created_at=1755077029, error=Error(code=None, message=None, param=N

```

▼ Monitoring the fine-tunes

```

1 import pandas as pd
2 from openai import OpenAI
3 client = OpenAI()
4 # Assume client is already set up and authenticated
5 response = client.fine_tuning.jobs.list(limit=3)# increase to see history
6
7 # Initialize lists to store the extracted data
8 job_ids = []
9 created_ats = []
10 statuses = []
11 models = []
12 training_files = []
13 error_messages = []
14 fine_tuned_models = [] # List to store the fine-tuned model names
15
16 # Iterate over the jobs in the response
17 for job in response.data:
18     job_ids.append(job.id)
19     created_ats.append(job.created_at)
20     statuses.append(job.status)
21     models.append(job.model)
22     training_files.append(job.training_file)
23     error_message = job.error.message if job.error else None
24     error_messages.append(error_message)
25
26     # Append the fine-tuned model name
27     fine_tuned_model = job.fine_tuned_model if hasattr(job, 'fine_tuned_model') else None
28     fine_tuned_models.append(fine_tuned_model)
29
30 # Create a DataFrame
31 df = pd.DataFrame({
32     'Job ID': job_ids,
33     'Created At': created_ats,
34     'Status': statuses,
35     'Model': models,
36     'Training File': training_files,
37     'Error Message': error_messages,
38     'Fine-Tuned Model': fine_tuned_models # Include the fine-tuned model names
39 })
40
41 # Convert timestamps to readable format
42 df['Created At'] = pd.to_datetime(df['Created At'], unit='s')
43 df = df.sort_values(by='Created At', ascending=False)
44
45 # Display the DataFrame
46 df

```

	Job ID	Created At	Status	Model	Training File	Error Message	Fine-Tuned Model
0	ftjob-f9JxcdrNVYIVnO4M3EZkLaLT	2025-08-13 09:23:49	succeeded	gpt-4.1-mini-2025-04-14	file-6VgWqp3wmd6Tos6aUihLLZ	None	ft:gpt-4.1-mini-2025-04-14:personal::C42Vkm29
1	ftjob-5BnCL6gUsFCTzyDYAN4CfvIZ	2024-12-02 10:04:26	succeeded	gpt-4o-mini-2024-07-18	file-KYWaDkNUDFSqqUcJnSv8hw	None	ft:gpt-4o-mini-2024-07-18:personal::AZy8s878
~	ftjob-	2024-12-02	...	gpt-4o-	file-	..	ft:gpt-4o-mini-2024-07-

✓ Make sure to obtain your fine-tune model here

If your OpenAI notifications are activated you should receive an email.

Otherwise run the "Monitoring the fine-tunes" cell above to check the status of your fine-tune job.

```

1 import pandas as pd
2
3 generation=False # False until the last model fine-tuned is found. Make sure it used the dataset you trained it on!
4 # Attempt to find the first non-empty Fine-Tuned Model
5 non_empty_models = df[df['Fine-Tuned Model'].notna() & (df['Fine-Tuned Model'] != '')]
6
7 if not non_empty_models.empty:
8     first_non_empty_model = non_empty_models['Fine-Tuned Model'].iloc[0]
9     print("The latest fine-tuned model is:", first_non_empty_model)
10    generation=True
11 else:
12     first_non_empty_model='None'
13     print("No fine-tuned models found.")

```

🔗 The latest fine-tuned model is: ft:gpt-4.1-mini-2025-04-14:personal::C42Vkm29

```

1 # Fine-tuned model found(True) or not(False)
2 generation

```

🔗 True

Note: Only continue to Step 3, to use the fine-tuned model when your fine-tuned model is ready. If your OpenAI notifications is activated, you will receive an email with the status of your fine-tuning job.

✓ 3.Using the fine-tuned OpenAI model

Note: There is a fine-tuning. As such, be patient! Run the `Monitoring the fine-tunes` cell and the `first_non_empty_model` cell from time to time.

If the fine-tuning succeeded and your model is ready, the name of your model will be `first_non_empty_model`

1. Go to the OpenAI Playground to test your model: <https://platform.openai.com/playground>

2. Check the metrics in the fine-tuning UI: <https://platform.openai.com/finetune/>

3. Try the fine-tuned model out in the cell below.

```

1 # Define the prompt
2 prompt="Which prize did Frederick Buechner create?"

```

Note: Only run the following cell if your fine-tune job has succeeded and a fine-tuned model is found in the *Monitoring the fine-tunes* section of *2. Fine-tuning the model*.

```

1 # Assume first_non_empty_model is defined above this snippet
2 if generation==True:
3     response = client.chat.completions.create(
4         # fine-tuned model or fallback if the model is not fine-tuned yet
5         model=first_non_empty_model or "gpt-4.1-mini-2025-04-14",
6         temperature=0.0, # Adjust as needed for variability
7         messages=[
8             {"role": "system", "content": "Given a question, reply with a complete explanation for students."},
9             {"role": "user", "content": prompt}
10        ]
11    )
12 else:
13     print("Error: Model is None, cannot proceed with the API request.")

```

```

1 if generation==True:
2     print(response)

```

➦ ChatCompletion(id='chatcmpl-C42hsGT0RiJhxJmjhV0kI7VCERHOX', choices=[Choice(finish_reason='stop', index=0, logprobs=None)], created=1744241100, model='gpt-4.1-mini-2025-04-14', object='chat.completion', usage={'completion_tokens': 12, 'prompt_tokens': 14, 'total_tokens': 26})

```

1 if (generation==True):
2     # Access the response from the first choice
3     response_text = response.choices[0].message.content
4     # Print the response
5     print(response_text)

```

➦ Frederick Buechner Prize for Preaching Explanation: The university is affiliated with the Congregation of Holy Cross (CSC), a Roman Catholic religious order. The university is affiliated with the Congregation of Holy Cross (CSC), a Roman Catholic religious order. The university is affiliated with the Congregation of Holy Cross (CSC), a Roman Catholic religious order.

```

1 import textwrap
2
3 if generation==True:
4     wrapped_text = textwrap.fill(response_text.strip(), 60)
5     print(wrapped_text)

```

➦ Frederick Buechner Prize for Preaching Explanation: The university is affiliated with the Congregation of Holy Cross (CSC), a Roman Catholic religious order. The university is affiliated with the Congregation of Holy Cross (CSC), a Roman Catholic religious order. The university is affiliated with the Congregation of Holy Cross (CSC), a Roman Catholic religious order. While religious affiliation is not a criterion for admission, more than 93% of students identify as Christian, with over 80% of the total being Catholic. Collectively, Catholic Mass is celebrated over 100 times per week on campus, and a large campus ministry program provides for the spiritual needs of the community. There are multitudes of religious statues and artwork around campus, most prominent of which are the statue of Mary on the Main Building, the Notre Dame Grotto, and the Word of Life mural on Hesburgh Library depicting Christ as a teacher. Additionally, every classroom displays a crucifix. There are many religious clubs (catholic and non-Catholic) at the university, including Council #1477 of the Knights of Columbus (KOC), Baptist Collegiate Ministry (BCM), Jewish Club, Muslim Student Association, Orthodox Christian Fellowship, The Mormon Club, and many more. The Notre Dame KofC are known for being the first collegiate council of KofC, operating a charitable concession stand during every home football game and owning their own building on campus which can be used as a cigar lounge. Fifty-seven chapels are located throughout the campus.

[Consult OpenAI fine-tune documentation for more](#)